

NETFLIX MOVIE RECOMMENDATION SYSTEM

By D Yakshitha(1RN18IS035)

1. Business Problem

1.1 Problem Description: Netflix is all about connecting people to the movies they love. To help customers find those movies, they developed world-class movie recommendation system: CinematchSM. Its job is to predict whether someone will enjoy a movie based on how much they liked or disliked other movies. Netflix use those predictions to make personal movie recommendations based on each customer's unique tastes. And while Cinematch is doing pretty well, it can always be made better.

Now there are a lot of interesting alternative approaches to how Cinematch works that netflix haven't tried. Some are described in the literature, some aren't. We're curious whether any of these can beat Cinematch by making better predictions. Because, frankly, if there is a much better approach it could make a big difference to our customers and our business.

1.2 Problem Statement:

Netflix provided a lot of rating data from which a recommendation system is built. Given different details about users and movies problem is to predict movies that would be best rated by a user

1.3 Sources : Data: <https://www.kaggle.com/rounakbanik/the-movies-dataset>
(<https://www.kaggle.com/rounakbanik/the-movies-dataset>) surprise library: <http://surpriselib.com/>
(<http://surpriselib.com/>) surprise library doc:
http://surprise.readthedocs.io/en/stable/getting_started.html
(http://surprise.readthedocs.io/en/stable/getting_started.html) (we use many models from this library) installing surprise: <https://github.com/NicolasHug/Surprise#installation>
(<https://github.com/NicolasHug/Surprise#installation>) Research paper:
<http://courses.ischool.berkeley.edu/i290-dm/s11/SECURE/a1-koren.pdf>
(<http://courses.ischool.berkeley.edu/i290-dm/s11/SECURE/a1-koren.pdf>) (most of our work was inspired by this paper) Best Algorithm(SVD):https://github.com/nishantml/NETFLIX-MOVIE-RECOMMENDATION-SYSTEM/blob/master/Netflix_Movie.ipynb
(https://github.com/nishantml/NETFLIX-MOVIE-RECOMMENDATION-SYSTEM/blob/master/Netflix_Movie.ipynb)

1.4 Real world/Business Objectives and constraints:

Objectives:

Predict the rating that a user would give to a movie that he has not yet rated. Minimize the difference between predicted and actual rating (RMSE and MAPE)

Constraints:

Some form of interpretability.

2. Machine Learning Problem

2.1 Data 2.1.1 Data Overview Get the data from : <https://www.kaggle.com/rounakbanik/the-movies-dataset> (<https://www.kaggle.com/rounakbanik/the-movies-dataset>) Data files : credits.csv keywords.csv links.csv links_small.csv movies_metadata.csv ratings.csv ratings_small.csv

*credits file contains cast and crew details of every movie. *keywords file contains Cast and Crew Information for all movies in the movies_metadata.csv file. *links file contains IMDB and TMDB IDs of all movies featured in the ratings.csv file (About 45,000 movies). *links_small file contains IMDB and TMDB IDs of all movies featured in the ratings_small.csv file (About 9000 movies).

*movies_metadata has 24 columns representing information about every movie like genres,id,production company,voting,run time,budget and so on. *ratings file contains ratings given by users to movies. *ratings_small file contains 100 ratings from 700 users on 9,000 +91 88708 42439. Is a subset of the ratings available in the Full MovieLens dataset.

2.2 Mapping the real world problem to a Machine Learning Problem:

2.2.1 Type of Machine Learning Problem For a given movie and user we need to predict the similar movies for which highest ratings would be given by him/her to the movie. The given problem is a Recommendation problem It can also seen as a Regression problem

2.2.2 Performance metric Mean Absolute Percentage Error:

https://en.wikipedia.org/wiki/Mean_absolute_percentage_error

(https://en.wikipedia.org/wiki/Mean_absolute_percentage_error) Root Mean Square Error:

https://en.wikipedia.org/wiki/Root-mean-square_deviation (https://en.wikipedia.org/wiki/Root-mean-square_deviation)

2.2.3 Machine Learning Objective and Constraints Minimize RMSE. Try to provide some interpretability.

3. Exploratory Data Analysis

3.1 Preprocessing

```
In [1]: from mpl_toolkits.mplot3d import Axes3D
from sklearn.preprocessing import StandardScaler
import matplotlib.pyplot as plt # plotting
import numpy as np # linear algebra
import os # accessing directory structure
import pandas as pd
```

```
In [2]: # Distribution graphs (histogram/bar graph) of column data
def plotPerColumnDistribution(df, nGraphShown, nGraphPerRow):
    nunique = df.nunique()
    df = df[[col for col in df if nunique[col] > 1 and nunique[col] < 50]] # For
    nRow, nCol = df.shape
    columnNames = list(df)
    nGraphRow = (nCol + nGraphPerRow - 1) // nGraphPerRow
    plt.figure(num = None, figsize = (6 * nGraphPerRow, 8 * nGraphRow), dpi = 80)
    for i in range(min(nCol, nGraphShown)):
        plt.subplot(nGraphRow, nGraphPerRow, i + 1)
        columnDf = df.iloc[:, i]
        if (not np.issubdtype(type(columnDf.iloc[0]), np.number)):
            valueCounts = columnDf.value_counts()
            valueCounts.plot.bar()
        else:
            columnDf.hist()
        plt.ylabel('counts')
        plt.xticks(rotation = 90)
        plt.title(f'{columnNames[i]} (column {i})')
    plt.tight_layout(pad = 1.0, w_pad = 1.0, h_pad = 1.0)
    plt.show()
```

```
In [3]: # Correlation matrix
def plotCorrelationMatrix(df, graphWidth):
    filename = df.dataframeName
    df = df.dropna('columns') # drop columns with NaN
    df = df[[col for col in df if df[col].nunique() > 1]] # keep columns where there
    if df.shape[1] < 2:
        print(f'No correlation plots shown: The number of non-NaN or constant columns')
        return
    corr = df.corr()
    plt.figure(num=None, figsize=(graphWidth, graphWidth), dpi=80, facecolor='w')
    corrMat = plt.matshow(corr, fignum = 1)
    plt.xticks(range(len(corr.columns)), corr.columns, rotation=90)
    plt.yticks(range(len(corr.columns)), corr.columns)
    plt.gca().xaxis.tick_bottom()
    plt.colorbar(corrMat)
    plt.title(f'Correlation Matrix for {filename}', fontsize=15)
    plt.show()
```

```
In [4]: # Scatter and density plots
def plotScatterMatrix(df, plotSize, textSize):
    df = df.select_dtypes(include=[np.number]) # keep only numerical columns
    # Remove rows and columns that would lead to df being singular
    df = df.dropna('columns')
    df = df[[col for col in df if df[col].nunique() > 1]] # keep columns where there is more than one value
    columnNames = list(df)
    if len(columnNames) > 10: # reduce the number of columns for matrix inversion to max 10
        columnNames = columnNames[:10]
    df = df[columnNames]
    ax = pd.plotting.scatter_matrix(df, alpha=0.75, figsize=[plotSize, plotSize])
    corrs = df.corr().values
    for i, j in zip(*plt.np.triu_indices_from(ax, k = 1)):
        ax[i, j].annotate('Corr. coef = %.3f' % corrs[i, j], (0.8, 0.2), xycoords='axesfraction',
                           textcoords='axesfraction')
    plt.suptitle('Scatter and Density Plot')
    plt.show()
```

Reading credits.csv file

```
In [5]: nRowsRead = 1000 # specify 'None' if want to read whole file
# credits.csv has 45476 rows in reality, but we are only loading/previewing the first 1000 rows
df1 = pd.read_csv('C:\\Users\\YAKSHITHA DONTI\\Downloads\\archive\\credits.csv')
df1.dataframeName = 'credits.csv'
nRow, nCol = df1.shape
print(f'There are {nRow} rows and {nCol} columns')
```

There are 1000 rows and 3 columns

```
In [6]: df1.head(5)
```

```
Out[6]:
```

	cast	crew	id
0	[{'cast_id': 14, 'character': 'Woody (voice)', ...	[{'credit_id': '52fe4284c3a36847f8024f49', 'de...	862
1	[{'cast_id': 1, 'character': 'Alan Parrish', '...	[{'credit_id': '52fe44bfc3a36847f80a7cd1', 'de...	8844
2	[{'cast_id': 2, 'character': 'Max Goldman', 'c...	[{'credit_id': '52fe466a9251416c75077a89', 'de...	15602
3	[{'cast_id': 1, 'character': 'Savannah Vannah...	[{'credit_id': '52fe44779251416c91011acb', 'de...	31357
4	[{'cast_id': 1, 'character': 'George Banks', '...	[{'credit_id': '52fe44959251416c75039ed7', 'de...	11862

```
In [7]: plotPerColumnDistribution(df1, 10, 5)
```

<Figure size 2400x512 with 0 Axes>

Reading keywords.csv file

```
In [9]: nRowsRead = 1000 # specify 'None' if want to read whole file
# keywords.csv has 46419 rows in reality, but we are only loading/previewing the
df2 = pd.read_csv('C:\\Users\\YAKSHITHA DONTI\\Downloads\\archive\\keywords.csv')
df2.dataframeName = 'keywords.csv'
nRow, nCol = df2.shape
print(f'There are {nRow} rows and {nCol} columns')
```

There are 1000 rows and 2 columns

```
In [10]: df2.head(5)
```

```
Out[10]:
```

	id	keywords
0	862	[{'id': 931, 'name': 'jealousy'}, {'id': 4290, ...
1	8844	[{'id': 10090, 'name': 'board game'}, {'id': 1...
2	15602	[{'id': 1495, 'name': 'fishing'}, {'id': 12392...
3	31357	[{'id': 818, 'name': 'based on novel'}, {'id': ...
4	11862	[{'id': 1009, 'name': 'baby'}, {'id': 1599, 'n...

```
In [11]: plotPerColumnDistribution(df2, 10, 5)
```

<Figure size 2400x512 with 0 Axes>

```
In [ ]: Reading links.csv file
```

```
In [12]: nRowsRead = 1000 # specify 'None' if want to read whole file
# links_small.csv has 9125 rows in reality, but we are only loading/previewing th
df3 = pd.read_csv('C:\\Users\\YAKSHITHA DONTI\\Downloads\\archive\\links_small.c
df3.dataframeName = 'links_small.csv'
nRow, nCol = df3.shape
print(f'There are {nRow} rows and {nCol} columns')
```

There are 1000 rows and 3 columns

```
In [13]: df3.head(5)
```

```
Out[13]:
```

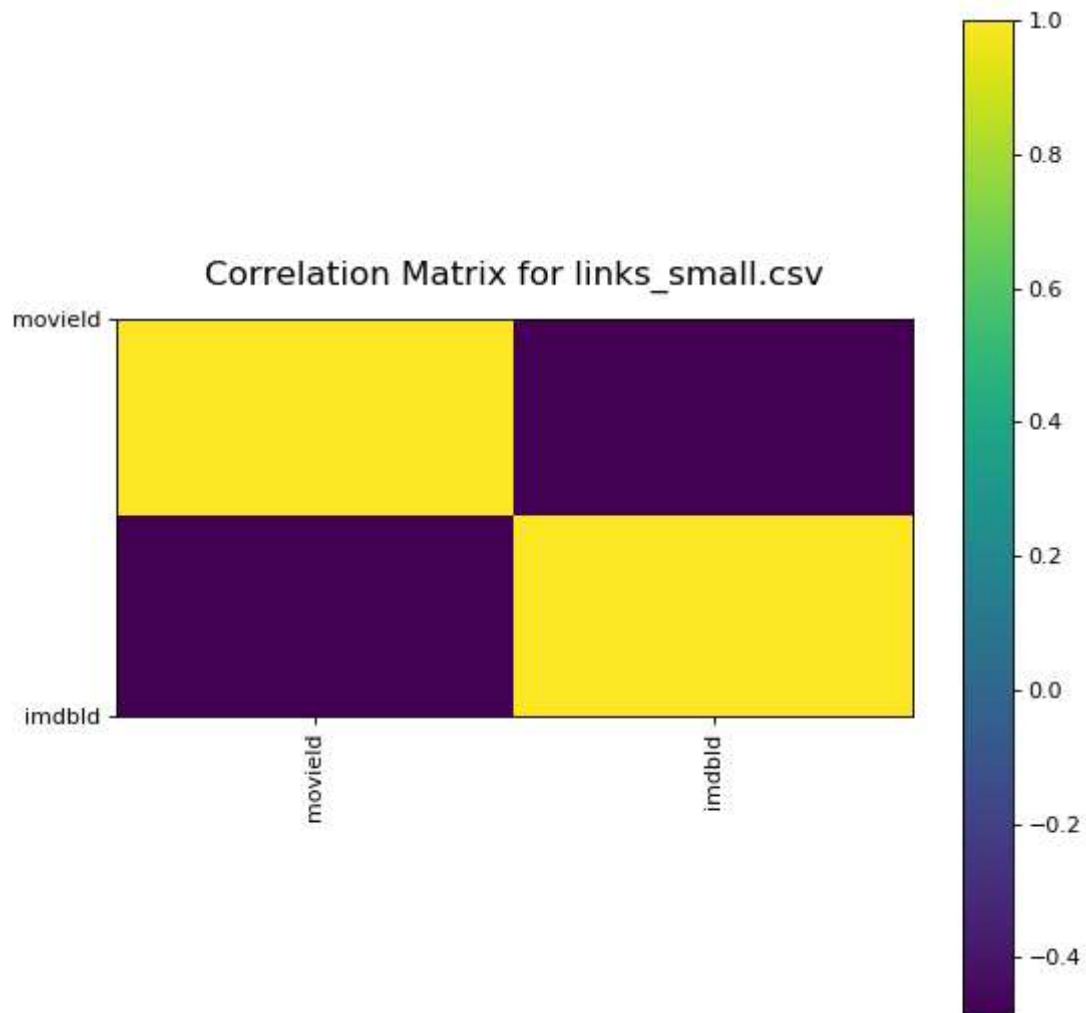
	movied	imdbid	tmdbid
0	1	114709	862.0
1	2	113497	8844.0
2	3	113228	15602.0
3	4	114885	31357.0
4	5	113041	11862.0

```
In [14]: plotPerColumnDistribution(df3, 10, 5)
```

<Figure size 2400x512 with 0 Axes>

Plotting heatmap between movieid and imdbid values

```
In [16]: plotCorrelationMatrix(df3, 8)
```




conclusion: movie id and imdb id are not correlated

Scatter and density plots for movie and imdb id

```
In [17]: plotScatterMatrix(df3, 9, 10)
```

Scatter and Density Plot



```
In [20]: %matplotlib inline
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from scipy import stats
from ast import literal_eval
from sklearn.feature_extraction.text import TfidfVectorizer, CountVectorizer
from sklearn.metrics.pairwise import linear_kernel, cosine_similarity
from nltk.stem.snowball import SnowballStemmer
from nltk.stem.wordnet import WordNetLemmatizer
from nltk.corpus import wordnet
from surprise import Reader, Dataset, SVD
from surprise.model_selection import cross_validate
import warnings; warnings.simplefilter('ignore')
```


Reading movies_metadata.csv file


```
In [21]: md = pd.read_csv('C:\\Users\\YAKSHITHA DONTI\\Downloads\\archive\\movies_metadata.csv')
md.head()
```


```
Out[21]:
```

	adult	belongs_to_collection	budget	genres	homepage	id	imdb_id
0	False	{'id': 10194, 'name': 'Toy Story Collection', ...}	300000000	[{'id': 16, 'name': 'Animation'}, {'id': 35, 'name': 'Family'}]	http://toystory.disney.com/toy-story	862	tt0114709
1	False	NaN	65000000	[{'id': 12, 'name': 'Adventure'}, {'id': 14, 'name': 'Fantasy'}]	NaN	8844	tt0113497
2	False	{'id': 119050, 'name': 'Grumpy Old Men Collect...	0	[{'id': 10749, 'name': 'Romance'}, {'id': 35, 'name': 'Family'}]	NaN	15602	tt0113228
3	False	NaN	16000000	[{'id': 35, 'name': 'Comedy'}, {'id': 18, 'name': 'Drama'}]	NaN	31357	tt0114885
4	False	{'id': 96871, 'name': 'Father of the Bride Col...	0	[{'id': 35, 'name': 'Comedy'}]	NaN	11862	tt0113041

5 rows × 24 columns



```
In [22]: md['genres'] = md['genres'].fillna('').apply(literal_eval).apply(lambda x: [i[0] for i in x])
md['year'] = pd.to_datetime(md['release_date'], errors='coerce').apply(lambda x: x.year if x else None)
```



```
In [23]: C = md['vote_average'].mean()
m = md['vote_count'].quantile(0.95)
print("C is %f, and m is %d"%(C,m))
```

C is 5.618207, and m is 434

```
In [24]: qualified = md[(md['vote_count'] >= m) & (md['vote_count'].notnull()) & (md['vote_average'] > C)]
qualified.shape
```

```
Out[24]: (2274, 25)
```

```
In [25]: def weighted_rating(x):
          v = x['vote_count']
          R = x['vote_average']
          return (v/(v+m) * R) + (m/(m+v) * C)
qualified = qualified[['title', 'year', 'vote_count', 'vote_average', 'popularity',
qualified['score'] = qualified.apply(weighted_rating, axis=1)
qualified = qualified.sort_values('score', ascending=False)
qualified.head(10)
```

```
Out[25]:
```

	title	year	vote_count	vote_average	popularity	genres	score
314	The Shawshank Redemption	1994	8358.0	8.5	51.6454	[Drama, Crime]	8.357746
834	The Godfather	1972	6024.0	8.5	41.1093	[Drama, Crime]	8.306334
12481	The Dark Knight	2008	12269.0	8.3	123.167	[Drama, Action, Crime, Thriller]	8.208376
2843	Fight Club	1999	9678.0	8.3	63.8696	[Drama]	8.184899
292	Pulp Fiction	1994	8670.0	8.3	140.95	[Thriller, Crime]	8.172155
351	Forrest Gump	1994	8147.0	8.2	48.3072	[Comedy, Drama, Romance]	8.069421
522	Schindler's List	1993	4436.0	8.3	41.7251	[Drama, History, War]	8.061007
23673	Whiplash	2014	4376.0	8.3	64.3	[Drama]	8.058025
5481	Spirited Away	2001	3968.0	8.3	41.0489	[Fantasy, Adventure, Animation, Family]	8.035598
1154	The Empire Strikes Back	1980	5998.0	8.2	19.471	[Adventure, Action, Science Fiction]	8.025793

Checking for rows containing NaN values

```
In [27]: links = pd.read_csv('C:\\Users\\YAKSHITHA DONTI\\Downloads\\archive\\links_small.csv')
links = links[links['tmdbId'].notnull()]['tmdbId'].astype('int')
print (md[pd.to_numeric(md['id'], errors='coerce').isnull()])
```

	adult	
19730		- Written by Ørnås
29503		Rune Balot goes to a casino connected to the ...
35587		Avalanche Sharks tells the story of a bikini ...

	belongs_to_collection		budget	
19730	0.065736	/ff9qCepilowshEtG2GYWwzt2bs4.jpg		
29503	1.931659	/zV8bHuSL6WXoD6FWogP9j4x80bL.jpg		
35587	2.185485	/zaSf50G7V8X8gqFvly88zDdRm46.jpg		

	genres	
19730	[Carousel Productions, Vision View Entertainme...	
29503	[Aniplex, GoHands, BROSTA TV, Mardock Scramble...	
35587	[Odyssey Media, Pulser Productions, Rogue Stat...	

	homepage		id	imdb_id	
19730	[{'iso_3166_1': 'CA', 'name': 'Canada'}, {'iso...	1997-08-20	0		
29503	[{'iso_3166_1': 'US', 'name': 'United States o...	2012-09-29	0		
35587	[{'iso_3166_1': 'CA', 'name': 'Canada'}]	2014-01-01	0		

	original_language		original_title	overview	
19730	104.0	[{'iso_639_1': 'en', 'name': 'English'}]	Released		
29503	68.0	[{'iso_639_1': 'ja', 'name': '日本語'}]	Released		
35587	82.0	[{'iso_639_1': 'en', 'name': 'English'}]	Released		

	...	revenue	runtime	spoken_languages	status	tagline	title	video	
19730	...	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
29503	...	NaN	NaN	NaN	NaN	NaN	NaN	NaN	
35587	...	NaN	NaN	NaN	NaN	NaN	NaN	NaN	

	vote_average	vote_count	year
19730	NaN	NaN	NaT
29503	NaN	NaN	NaT
35587	NaN	NaN	NaT

[3 rows x 25 columns]



```
In [28]: md = md.drop([19730, 29503, 35587])
md['id'] = md['id'].astype('int')
```

Getting recommendations based on a movie

```
In [29]: def get_recommendations(title, cosine_sim):  
    idx = indices[title]  
    sim_scores = list(enumerate(cosine_sim[idx]))  
    sim_scores = sorted(sim_scores, key=lambda x: x[1], reverse=True)  
    sim_scores = sim_scores[1:31]  
    movie_indices = [i[0] for i in sim_scores]  
    return titles.iloc[movie_indices]
```

Merging credits and keywords file

```
In [31]: credits = pd.read_csv('C:\\Users\\YAKSHITHA DONTI\\Downloads\\archive\\credits.csv')  
keywords = pd.read_csv('C:\\Users\\YAKSHITHA DONTI\\Downloads\\archive\\keywords.csv')  
keywords['id'] = keywords['id'].astype('int')  
credits['id'] = credits['id'].astype('int')  
md = md.merge(credits, on='id')  
md = md.merge(keywords, on='id')  
smd1 = md[md['id'].isin(links)]  
  
features = ['cast', 'crew', 'keywords']  
for feature in features:  
    smd1[feature] = smd1[feature].apply(literal_eval)
```

```
In [32]: def get_director(x):
        for i in x:
            if i['job'] == 'Director':
                return i['name']
        return np.nan
smd1['director'] = smd1['crew'].apply(get_director)
smd1.head()
```

Out[32]:

overview	...	tagline	title	video	vote_average	vote_count	year	cast
Led by Woody, Andy's toys live happily in his	NaN	Toy Story	False	7.7	5415.0	1995	{'cast_id': 14, 'character': 'Woody (voice)',...
When siblings Judy and Peter discover an encha...	...	Roll the dice and unleash the excitement!	Jumanji	False	6.9	2413.0	1995	{'cast_id': 1, 'character': 'Alan Parrish', '...
A family wedding reignites the ancient feud be...	...	Still Yelling. Still Fighting. Still Ready for...	Grumpier Old Men	False	6.5	92.0	1995	{'cast_id': 2, 'character': 'Max Goldman', 'c...
Cheated on, mistreated and stepped on, the wom...	...	Friends are the people who let you be yourself...	Waiting to Exhale	False	6.1	34.0	1995	{'cast_id': 1, 'character': 'Savannah Vannah...
ust when George banks has ecovered om his	Just When His World Is Back To Normal... He's ...	Father of the Bride Part II	False	5.7	173.0	1995	{'cast_id': 1, 'character': 'George Banks', '...



```
In [33]: def get_list(x):
        if isinstance(x, list):
            names = [i['name'] for i in x]
            #Check if more than 3 elements exist. If yes, return only first three. If
            if len(names) > 3:
                names = names[:3]
            return names
        return []
features = ['cast', 'keywords']
for feature in features:
    smd1[feature] = smd1[feature].apply(get_list)
smd1[['title', 'cast', 'director', 'keywords', 'genres']].head(3)
```

Out[33]:

	title	cast	director	keywords	genres
0	Toy Story	[Tom Hanks, Tim Allen, Don Rickles]	John Lasseter	[jealousy, toy, boy]	[Animation, Comedy, Family]
1	Jumanji	[Robin Williams, Jonathan Hyde, Kirsten Dunst]	Joe Johnston	[board game, disappearance, based on children'...	[Adventure, Fantasy, Family]
2	Grumpier Old Men	[Walter Matthau, Jack Lemmon, Ann-Margret]	Howard Deutch	[fishing, best friend, duringcreditsstinger]	[Romance, Comedy]


```
In [34]: def clean_data(x):
    if isinstance(x, list):
        return [str.lower(i.replace(" ", "")) for i in x]
    else:
        #Check if director exists. If not, return empty string
        if isinstance(x, str):
            return str.lower(x.replace(" ", ""))
        else:
            return ''
# Apply clean_data function to your features.
features = ['cast', 'keywords', 'director', 'genres']

for feature in features:
    smd1[feature] = smd1[feature].apply(clean_data)
smd1['director'] = smd1['director'].apply(lambda x: [x,x, x])
smd1.head(3)
```

Out[34]:

	adult	belongs_to_collection	budget	genres	homepage	id	imdb_id
0	False	{'id': 10194, 'name': 'Toy Story Collection', ...}	30000000	[animation, comedy, family]	http://toystory.disney.com/toy-story	862	tt0114709
1	False	NaN	65000000	[adventure, fantasy, family]	NaN	8844	tt0113497
2	False	{'id': 119050, 'name': 'Grumpy Old Men Collect...	0	[romance, comedy]	NaN	15602	tt0113228

3 rows × 29 columns



```
In [35]: def create_soup(x):
          return ' '.join(x['keywords']) + ' ' + ' '.join(x['cast']) + ' '.join(x['director'])
smd1['soup'] = smd1.apply(create_soup, axis=1)
smd1[['title', 'cast', 'director', 'keywords', 'genres', 'soup']].head(3)
```

```
Out[35]:
```

	title	cast	director	keywords	genres	soup
0	Toy Story	[tomhanks, timallen, donrickles]	[johnlasseter, johnlasseter, johnlasseter]	[jealousy, toy, boy]	[animation, comedy, family]	jealousy toy boy tomhanks timallen donricklesj...
1	Jumanji	[robinwilliams, jonathanhyde, kirstendunst]	[joejohnston, joejohnston, joejohnston]	[boardgame, disappearance, basedonchildren'sbook]	[adventure, fantasy, family]	boardgame disappearance basedonchildren'sbook ...
2	Grumpier Old Men	[waltermatthau, jacklemmon, ann-margret]	[howarddeutch, howarddeutch, howarddeutch]	[fishing, bestfriend, duringcreditsstinger]	[romance, comedy]	fishing bestfriend duringcreditsstinger walter...

4. Machine Learning Model

4.1 Using countvectorizer and cosine similarity to get the similar movies that would be liked by user

```
In [57]: count = CountVectorizer(analyzer='word', ngram_range=(1, 2), min_df=0, stop_words=
count_matrix = count.fit_transform(smd1['soup'])

cosine_sim = cosine_similarity(count_matrix, count_matrix)

smd1 = smd1.reset_index()
titles = smd1['title']
indices = pd.Series(smd1.index, index=smd1['title'])

indices.head()
```

```
Out[57]: title
Toy Story      0
Jumanji        1
Grumpier Old Men  2
Waiting to Exhale  3
Father of the Bride Part II  4
dtype: int64
```

Getting similar movies to "Toy Story" with cosine similarity


```
In [37]: get_recommendations('Toy Story',cosine_sim)
```

```
Out[37]: 2522 Toy Story 2
7629 Toy Story 3
8519 Toy Story of Terror!
6496 Cars
7914 Cars 2
6386 Luxo Jr.
283 The Santa Clause
320 The Flintstones
6534 Monster House
1883 A Bug's Life
246 Man of the House
1432 Meet the Deedles
1593 Freaky Friday
2751 Creature Comforts
4341 The Looney, Looney, Looney Bugs Bunny Movie
7355 Aliens in the Attic
584 Operation Dumbo Drop
933 The Wrong Trousers
1633 Hot Lead & Cold Feet
3333 See Spot Run
5287 Bon Voyage, Charlie Brown (and Don't Come Back!)
7541 The Spy Next Door
8040 We Bought a Zoo
3019 The Adventures of Rocky & Bullwinkle
749 The Little Princess
2194 Inspector Gadget
4246 Captain Ron
9017 Freaky Friday
81 Dunston Checks In
137 Casper
Name: title, dtype: object
```

```
In [39]: reader = Reader()
ratings = pd.read_csv('C:\\Users\\YAKSHITHA DONTI\\Downloads\\archive\\ratings_1.csv')
ratings.head()
```

```
Out[39]:
```

	userId	movieId	rating	timestamp
0	1	31	2.5	1260759144
1	1	1029	3.0	1260759179
2	1	1061	3.0	1260759182
3	1	1129	2.0	1260759185
4	1	1172	4.0	1260759205

5. Cross Validation (Performance measurement RMSE and MAE)

```
In [41]: from surprise.model_selection import KFold
data = Dataset.load_from_df(ratings[['userId', 'movieId', 'rating']], reader)
kf=KFold(n_splits=5)
kf.split(data)
svd = SVD()
cross_validate(svd, data, measures=['RMSE', 'MAE'])
```

```
Out[41]: {'test_rmse': array([0.90161518, 0.90036053, 0.89069071, 0.90108111, 0.8939474
]),
'test_mae': array([0.6944839 , 0.69020138, 0.6867198 , 0.69153325, 0.6908901
4]),
'fit_time': (3.2475314140319824,
3.219769239425659,
3.2632484436035156,
3.2235045433044434,
3.2367424964904785),
'test_time': (0.19293832778930664,
0.08161067962646484,
0.1631002426147461,
0.08220863342285156,
0.08162975311279297)}
```

rmse and mae values for 5 splits cross validation

```
In [42]: from surprise import accuracy
from surprise.model_selection import train_test_split
trainset = data.build_full_trainset()
algo=SVD()
algo.fit(trainset)
ratings[ratings['userId'] == 1]
```

```
Out[42]:
```

	userId	movieId	rating	timestamp
0	1	31	2.5	1260759144
1	1	1029	3.0	1260759179
2	1	1061	3.0	1260759182
3	1	1129	2.0	1260759185
4	1	1172	4.0	1260759205
5	1	1263	2.0	1260759151
6	1	1287	2.0	1260759187
7	1	1293	2.0	1260759148
8	1	1339	3.5	1260759125
9	1	1343	2.0	1260759131
10	1	1371	2.5	1260759135
11	1	1405	1.0	1260759203
12	1	1953	4.0	1260759191
13	1	2105	4.0	1260759139
14	1	2150	3.0	1260759194
15	1	2193	2.0	1260759198
16	1	2294	2.0	1260759108
17	1	2455	2.5	1260759113
18	1	2968	1.0	1260759200
19	1	3671	3.0	1260759117

List of user id 1 ratings

```
In [43]: svd.predict(1, 302)
```

```
Out[43]: Prediction(uid=1, iid=302, r_ui=None, est=3.0539590474992013, details={'was_impossible': False})
```

It says user id 1 would give 3.05 as rating to movie id 302

```
In [50]: def convert_int(x):
          try:
              return int(x)
          except:
              return np.nan

id_map = pd.read_csv('C:\\Users\\YAKSHITHA DONTI\\Downloads\\archive\\links_small.csv')
id_map['tmdbId'] = id_map['tmdbId'].apply(convert_int)
id_map.columns = ['movieId', 'id']
id_map = id_map.merge(smd1[['title', 'id']], on='id').set_index('title')
indices_map = id_map.set_index('id')
```

6. Hybrid Model

6.1 Making a hybrid model to predict the user rating. Cosine similarity gets similar movies to passed movie and SVD will predict ratings given by the user to those similar movies and recommends top 10 movies to the user!!!

And it's done.....That's recommendation system!

```
In [63]: def hybrid(userId, title):
    idx = indices[title]
    tmdbId = id_map.loc[title]['id']
    #print(idx)
    movie_id = id_map.loc[title]['movieId']

    sim_scores = list(enumerate(cosine_sim[int(idx)]))
    sim_scores = sorted(sim_scores, key=lambda x: x[1], reverse=True)
    sim_scores = sim_scores[1:26]
    movie_indices = [i[0] for i in sim_scores]
    movies = smd1.iloc[movie_indices][['title', 'vote_count', 'vote_average', 'year']]
    movies['est'] = movies['id'].apply(lambda x: svd.predict(userId, indices_map.get(x)))
    movies = movies.sort_values('est', ascending=False)
    return movies.head(10)
hybrid(1, 'Avatar')
#hybrid(500, 'Avatar')
```

```
Out[63]:
```

	title	vote_count	vote_average	year	id	est
1011	The Terminator	4208.0	7.4	1984	218	3.360299
522	Terminator 2: Judgment Day	4274.0	7.7	1991	280	3.026912
8658	X-Men: Days of Future Past	6155.0	7.5	2014	127585	2.998357
9004	Suicide Squad	7717.0	5.9	2016	297761	2.983292
3181	The Time Machine	217.0	7.5	1960	2134	2.970603
2131	Superman	1042.0	6.9	1978	1924	2.966292
3060	Sinbad and the Eye of the Tiger	39.0	6.3	1977	11940	2.952807
5310	Frank Herbert's Dune	114.0	6.7	2000	876	2.927986
2396	Time Bandits	255.0	6.6	1981	36819	2.918542
5559	Slaughterhouse-Five	47.0	6.3	1972	24559	2.898039

Type *Markdown* and LaTeX: α^2